

lab7

October 11, 2016

1 EE 379K: Lab 7

1.1 Rohan Nagar and Wenyang Fu

```
In [1]: import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn import linear_model
from sklearn import cross_validation
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.grid_search import GridSearchCV

from statsmodels.formula.api import ols

%matplotlib inline
```

```
C:\Anaconda3\lib\site-packages\sklearn\cross_validation.py:44: DeprecationWarning: This module was deprecated in
"0.18.0; it will be removed in 0.20.0.", DeprecationWarning)
C:\Anaconda3\lib\site-packages\sklearn\grid_search.py:43: DeprecationWarning: This module was deprecated in
"0.18.0; it will be removed in 0.20.0.", DeprecationWarning)
```

2 Problem 1

This question should be answered using the `Carseats` data set.

```
In [2]: carseats = pd.read_csv('Data/Carseats.csv')
carseats.drop(carseats.columns[0], axis=1, inplace=True)
carseats.head()
```

```
Out[2]:
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	\
0	9.50	138	73	11	276	120	Bad	42	
1	11.22	111	48	16	260	83	Good	65	
2	10.06	113	35	10	269	80	Medium	59	
3	7.40	117	100	4	466	97	Medium	55	
4	4.15	141	64	3	340	128	Bad	38	

	Education	Urban	US
0	17	Yes	Yes
1	10	Yes	Yes
2	12	Yes	Yes
3	14	Yes	Yes
4	13	Yes	No

2.0.1 Part A

Fit a multiple regression model to predict Sales using Price, Urban, and US.

```
In [3]: carseats = carseats.replace(['Yes', 'No'], [1, 0])

regr = linear_model.LinearRegression()
X = carseats[['Price', 'Urban', 'US']]
y = carseats['Sales']
regr.fit(X, y)

print('Coefficients: {}'.format(regr.coef_))
print('Intercept: {}'.format(regr.intercept_))
print('R^2: {}'.format(regr.score(X, y)))

model = ols("Sales ~ Price + Urban + US", carseats).fit()
model.summary()
```

```
Coefficients: [-0.05445885 -0.02191615  1.2005727 ]
Intercept: 13.043468936764896
R^2: 0.23927539218405547
```

```
Out[3]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                  0.239
Model:                            OLS    Adj. R-squared:             0.234
Method:                 Least Squares    F-statistic:                 41.52
Date:                  Tue, 11 Oct 2016    Prob (F-statistic):          2.39e-23
Time:                  22:05:41    Log-Likelihood:              -927.66
No. Observations:                400    AIC:                        1863.
Df Residuals:                   396    BIC:                        1879.
Df Model:                        3
Covariance Type:                nonrobust
=====
               coef    std err          t      P>|t|      [95.0% Conf. Int.]
-----
Intercept    13.0435     0.651     20.036     0.000     11.764     14.323
Price        -0.0545     0.005    -10.389     0.000     -0.065     -0.044
Urban        -0.0219     0.272     -0.081     0.936     -0.556     0.512
US           1.2006     0.259      4.635     0.000      0.691     1.710
=====
Omnibus:                 0.676    Durbin-Watson:              1.912
Prob(Omnibus):            0.713    Jarque-Bera (JB):            0.758
Skew:                     0.093    Prob(JB):                    0.684
Kurtosis:                 2.897    Cond. No.                     628.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""
```

2.0.2 Part B

Provide an interpretation of each coefficient in the model. Be careful - some of the variables in the model are qualitative!

2.0.3 Answer

1. The first coefficient denotes that more expensive car seats are less likely to be sold.
2. The second coefficient denotes that a rural car buyer is slightly more likely to purchase a car seat.
3. The third coefficient denotes that a car buyer in the US is far more likely to purchase a car seat than buyers outside the US.

2.0.4 Part C

Write out the model in equation form, being careful to handle the qualitative variables properly.

$$\text{Sales} = 13.0435 - .0544x_0 - .02192x_1 + 1.2005x_2$$

2.0.5 Part D

For which of the predictors can you reject the null hypothesis $H_0 : \beta_j = 0$?

2.0.6 Answer

Price and US, based on the p-values.

2.0.7 Part E

On the basis of your response to the previous question, fit a smaller model that only uses the predictors for which there is evidence of association with the outcome.

```
In [4]: regr = linear_model.LinearRegression()
        X = carseats[['Price', 'US']]
        y = carseats['Sales']
        regr.fit(X, y)

        print('Coefficients: {}'.format(regr.coef_))
        print('Intercept: {}'.format(regr.intercept_))
        print('R^2: {}'.format(regr.score(X, y)))

        model = ols("Sales ~ Price + US", carseats).fit()
        model.summary()
```

```
Coefficients: [-0.05447763  1.19964294]
Intercept: 13.03079275461576
R^2: 0.23926288842678567
```

```
Out[4]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Sales    R-squared:                0.239
Model:                            OLS    Adj. R-squared:           0.235
Method:                           Least Squares    F-statistic:             62.43
Date:                            Tue, 11 Oct 2016    Prob (F-statistic):       2.66e-24
Time:                            22:05:41    Log-Likelihood:          -927.66
No. Observations:                  400    AIC:                     1861.
Df Residuals:                      397    BIC:                     1873.
```

```

Df Model:                2
Covariance Type:         nonrobust
=====
              coef      std err          t      P>|t|      [95.0% Conf. Int.]
-----
Intercept    13.0308      0.631     20.652     0.000      11.790      14.271
Price       -0.0545      0.005    -10.416     0.000      -0.065      -0.044
US           1.1996      0.258      4.641     0.000       0.692       1.708
=====
Omnibus:                0.666   Durbin-Watson:                1.912
Prob(Omnibus):          0.717   Jarque-Bera (JB):                0.749
Skew:                   0.092   Prob(JB):                  0.688
Kurtosis:               2.895   Cond. No.                  607.
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""

```

2.0.8 Part F

How well do the models in (a) and (e) fit the data?

2.0.9 Answer

The two models fit similarly, with the model in (e) doing ever so slightly better.

2.0.10 Part G

Using the model from (e), obtain 95% confidence intervals for the coefficients.

2.0.11 Answer

1. Price: [-0.065, -0.044]
2. US: [0.692, 1.708]
3. Intercept: [11.790, 14.271]

2.0.12 Part H

Is there evidence of outliers or high leverage observations in the model from (e)?

2.0.13 Answer


No

```

In [5]: fig, ax = plt.subplots()
        len(regr.predict(X))
        # ax.scatter(regr.predict(X), regr.residues_)

```

```
Out[5]: 400
```



3 Problem 2

This problem involves the `Boston` data set, which we saw in the lab for this chapter. We will now try to predict per capita crime rate using the other variables in this data set. In other words, per capita crime rate is the response, and the other variables are the predictors.

```
In [6]: boston = pd.read_csv('Data/Boston.csv')
        boston.dropna()
        boston.head()
```

```
Out[6]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio \
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7

	black	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

3.0.1 Part A

For each predictor, fit a simple linear regression model to predict the response. Describe your results. In which of the models is there a statistically significant associate between the predictor and the response? Create some plots to back up your assertions.

3.0.2 Answer

Everything except *chas* is statistically significant.

```
In [7]: X = boston.drop(boston.columns[0], axis=1)
        y = boston.crim
        cols = X.columns
        scores = []
        coefficients = []
        for col in cols:
            regr = linear_model.LinearRegression()
            X_train = X[col].reshape(-1, 1)
            regr.fit(X_train, y)
            coefficients.append(regr.coef_)
            scores.append((col, regr.score(X_train, y)))
        print("Feature {}'s R^2 is: {}".format(col, regr.score(X_train, y)))
        # print('Coefficients: {}'.format(regr.coef_))
        # print('Intercept: {}'.format(regr.intercept_))
        print(sorted(scores, key=lambda x: x[1], reverse=True))
```

Feature zn's R² is: 0.04018790803211081

Feature indus's R² is: 0.16531007043075163

Feature chas's R² is: 0.0031238689633057426

Feature nox's R² is: 0.17721718179269375

Feature rm's R² is: 0.048069116716083604

Feature age's R² is: 0.12442145175894635

Feature dis's R² is: 0.1441493749253987

Feature rad's R² is: 0.39125668674998915

Feature tax's R² is: 0.3396142433788123

Feature ptratio's R² is: 0.0840684389437365

Feature black's R² is: 0.1482742394241312

Feature lstat's R² is: 0.20759093253433558

Feature medv's R² is: 0.15078046904975717

[('rad', 0.39125668674998915), ('tax', 0.3396142433788123), ('lstat', 0.20759093253433558), ('nox', 0.17721718179269375), ('indus', 0.16531007043075163), ('age', 0.12442145175894635), ('dis', 0.1441493749253987), ('black', 0.1482742394241312), ('zn', 0.04018790803211081), ('rm', 0.048069116716083604), ('ptratio', 0.0840684389437365), ('chas', 0.0031238689633057426)]

```
In [8]: sns.pairplot(boston, x_vars=X.columns, y_vars=['crim'])
```

Out[8]: <seaborn.axisgrid.PairGrid at 0x20bf7fdf400>



3.0.3 Part B

Fit a multiple regression model to predict the response using all of the predictors. Describe your results. For which predictors can we reject the null hypothesis $H_0 : \beta_j = 0$?

3.0.4 Answer

zn, dis, rad, black, medv

```
In [9]: all_columns = "+".join(boston.columns.difference(["crim"]))
        formula = "crim ~" + all_columns
```

```

model = ols(formula, boston).fit()
model.summary()

regr = linear_model.LinearRegression()
regr.fit(X, y)

```

Out[9]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

3.0.5 Part C

How do your results from (a) compare to your results from (b)? Create a plot displaying the univariate regression coefficients from (a) on the x-axis, and the multiple regression coefficients from (b) on the y-axis. That is, each predictor is displayed as a single point in the plot. Its coefficient in a simple linear regression model is shown on the x-axis, and its coefficient estimate in the multiple linear regression model is shown on the y-axis.

3.0.6 Answer


Coefficients are fairly correlated with the sole exception of *nox*, which is an outlier.

```

In [10]: fig, ax = plt.subplots()
         ax.scatter(coefficients, regr.coef_)

```

Out[10]: <matplotlib.collections.PathCollection at 0x20bfe8ab5c0>



3.0.7 Part D

Is there evidence of non-linear associate between any of the predictors and the response? To answer this question, for each predictor X , fit a model of the form

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon$$

Yes, using some R code, there was some non-linear association for all features except black and chas.

In []:

4 Problem 3

We will now try to predict per capita crime rate in the `Boston` data set.

4.0.1 Part A

Try out some of the regression methods explored in this chapter, such as the best subset selection, the lasso, ridge regression, and PCR. Present and discuss results for the approaches that you consider.

4.0.2 Approach 1: Lasso

Lasso does pretty well with a $R^2 = .538$. As you can see, it ends up shrinking the importance of 5 features, and retains the rest.

```
In [27]: X = boston.drop(boston.columns[0], axis=1)
        y = boston.crim
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)

        regr = linear_model.LassoCV(cv=5)
        regr.fit(X_train, y_train)
        print(regr.score(X_test, y_test))
        regr.coef_
```

0.537588283341

```
Out[27]: array([ 2.24761117e-02, -0.00000000e+00, -0.00000000e+00,
                -0.00000000e+00,  0.00000000e+00,  2.23614202e-02,
                -1.83611966e-02,  5.12322719e-01,  4.26053805e-05,
                -0.00000000e+00, -7.46120467e-03,  2.83184483e-02,
                -1.21316678e-01])
```

4.0.3 Approach 2: Ridge

Ridge does slightly better than lasso with a $R^2 = .557$. Unlike LASSO, it doesn't aggressively zero out as many features, but it does scatter the feature weights exponentially across four orders of magnitude.

```
In [28]: regr = linear_model.RidgeCV(cv=5)
        regr.fit(X_train, y_train)
        print(regr.score(X_test, y_test))
        regr.coef_
```

0.557024626663

```
Out[28]: array([ 4.36165878e-02, -3.27519656e-02, -9.88111269e-01,
                -1.08919102e+01,  5.24808970e-01,  1.00542481e-02,
                -9.92226506e-01,  6.20592963e-01, -3.75232375e-03,
                -2.92647312e-01, -5.21585510e-03,  3.22843876e-02,
                -2.31279174e-01])
```


4.0.4 Approach 3: Principal Components Regression (via Partial Least Squares)

PCR seems to do the worst with a $R^2 = .236$. By embedding the original data matrix onto a lower-dimensional subspace, PCR seems to have left out some important information.

```
In [29]: from sklearn.cross_decomposition import PLSRegression
```

```
regr = PLSRegression(n_components=2)
regr.fit(X_train, y_train)
print(regr.score(X_test, y_test))
regr.coef_
```

```
0.235949756809
```

```
Out[29]: array([[ 0.0283097 ],
 [ 0.01922917],
 [-1.54016339],
 [ 1.70274105],
 [ 0.33906391],
 [ 0.00185609],
 [-0.12586492],
 [ 0.26927111],
 [ 0.01053107],
 [ 0.19701884],
 [-0.01176867],
 [ 0.07584012],
 [-0.07348777]])
```

4.0.5 Part B

Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure that you are evaluating model performance using validation set error, cross-validation, or some other reasonable alternative, as opposed to using training error.

4.0.6 Proposal: Ridge Regression

Ridge regression seems to do well without any tuning, and it has built-in feature selection (so to speak), so we decided to tune it further. The best model selected by Grid Search performed slightly better than the base RidgeCV model we selected.

```
In [54]: from sklearn.model_selection import GridSearchCV
```

```
alphas = list(np.power(10., np.arange(-2, 5)))

param_grid = [
    {'alpha': alphas, 'normalize': [True, False]},
]

model = linear_model.RidgeCV(alphas, cv=5)
model.fit(X_train, y_train)
print(model.alpha_)
print(model.score(X_test, y_test))

greg = GridSearchCV(linear_model.Ridge(), param_grid=param_grid, cv=5)
greg.fit(X_train, y_train)
print(greg.best_params_)
```

```

print(greg.best_score_ )
print(greg.score(X_test, y_test))
print(greg.bestr_)

1000.0
0.556817916986
{'alpha': 0.10000000000000001, 'normalize': True}
0.463140350198
0.559753705589
GridSearchCV(cv=5, error_score='raise',
             estimator=Ridge(alpha=1.0, copy_X=True, fit_intercept=True, max_iter=None,
                             normalize=False, random_state=None, solver='auto', tol=0.001),
             fit_params={}, iid=True, n_jobs=1,
             param_grid=[{'alpha': [0.01, 0.10000000000000001, 1.0, 10.0, 100.0, 1000.0, 10000.0]}, 'normalize':
             pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
             scoring=None, verbose=0)

```

4.0.7 Part C

Does your chosen model involve all of the features in the data set? Why or why not?

4.0.8 Answer

Yes, our model did involve all of the features in the dataset. All of the features were deemed important enough by Ridge Regression (when $\alpha = 1000$) that they each had a final impact on the target variable.

In [40]: `print(model.coef_)`

```

[ 0.02674592 -0.06658358 -0.02817974 -0.00879101  0.03970368  0.01769596
 -0.24646743  0.47776279  0.00222968 -0.07355824 -0.0071971  0.05144703
 -0.1390575 ]

```

5 Problem 4

In this exercise, we will predict the number of applications received using the other variables in the College data set.

In [15]: `college = pd.read_csv('Data/College.csv')`
`college.head()`

```

Out[15]:
      Unnamed: 0 Private  Apps  Accept  Enroll  Top10perc  \
0  Abilene Christian University    Yes  1660    1232    721      23
1           Adelphi University    Yes  2186    1924    512      16
2           Adrian College      Yes  1428    1097    336      22
3       Agnes Scott College      Yes   417     349    137      60
4   Alaska Pacific University    Yes   193     146     55      16

      Top25perc  F.Undergrad  P.Undergrad  Outstate  Room.Board  Books  Personal  \
0           52          2885           537      7440        3300    450      2200
1           29          2683          1227     12280        6450    750      1500
2           50          1036            99     11250        3750    400      1165
3           89           510            63     12960        5450    450       875
4           44           249           869      7560        4120    800      1500

      PhD  Terminal  S.F.Ratio  perc.alumni  Expend  Grad.Rate

```

0	70	78	18.1	12	7041	60
1	29	30	12.2	16	10527	56
2	53	66	12.9	30	8735	54
3	92	97	7.7	37	19016	59
4	76	72	11.9	2	10922	15

5.0.1 Part A

Split the data set into a training set and a test set.

```
In [16]: college = college.replace(['Yes', 'No'], [1, 0])
```

```
In [17]: features = ['Private', 'Accept', 'Enroll', 'Top10perc', 'F.Undergrad',
                    'P.Undergrad', 'Outstate', 'Room.Board', 'Books', 'Personal',
                    'PhD', 'Terminal', 'S.F.Ratio', 'perc.alumni', 'Expend', 'Grad.Rate']
X_train, X_test, Y_train, Y_test = train_test_split(college[features], college['Apps'], test_s
```

5.0.2 Part B

Fit a linear model using least squares on the training set, and report the test error obtained.

```
In [18]: from sklearn.metrics import mean_squared_error
```

```
reg = linear_model.LinearRegression()
reg.fit(X_train, Y_train)

score = mean_squared_error(Y_test, reg.predict(X_test))
print("Error on test data: {}".format(score))
```

Error on test data: 1793667.3647683186

5.0.3 Part C

Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test error obtained.

```
In [19]: ridge_reg = linear_model.RidgeCV(alphas=[0.01, 0.1, 0.5, 1, 1.5, 2, 5, 10])
ridge_reg.fit(X_train, Y_train)

score = mean_squared_error(Y_test, ridge_reg.predict(X_test))
print("Error on test data: {}".format(score))
```

Error on test data: 1790804.8706922578

5.0.4 Part D

Fit a lasso model on the training set, with λ chosen by cross-validation. Report the test error obtained, along with the number of non-zero coefficient estimates.

```
In [20]: lasso_reg = linear_model.LassoCV(alphas=[0.01, 0.1, 0.5, 1, 1.5, 2, 5, 10])
lasso_reg.fit(X_train, Y_train)

score = mean_squared_error(Y_test, lasso_reg.predict(X_test))
print("Error on test data: {}".format(score))
```

Error on test data: 1793660.783999553

5.0.5 Part E

Fit a PCR model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
In [21]: from sklearn.decomposition import PCA
```

```
pca = PCA()
X_reduced = pca.fit_transform(X_train)

# Show how much variance is explained
np.cumsum(np.round(pca.explained_variance_ratio_, decimals=4)*100)
```

```
Out[21]: array([ 49.94,  87.47,  95.73,  97.5 ,  98.71,  99.45,  99.92,
                99.98, 100.01, 100.01, 100.01, 100.01, 100.01, 100.01,
                100.01, 100.01])
```

```
In [22]: n = len(X_reduced)
kf_10 = cross_validation.KFold(n, n_folds=10, shuffle=True, random_state=2)
```

```
# Use Linear Regression with increasing number of principal components
regr = linear_model.LinearRegression()
mse = []
```

```
# Without any components
score = -1*cross_validation.cross_val_score(regr, np.ones((n,1)), Y_train.ravel(), cv=kf_10, s
mse.append(score)
```

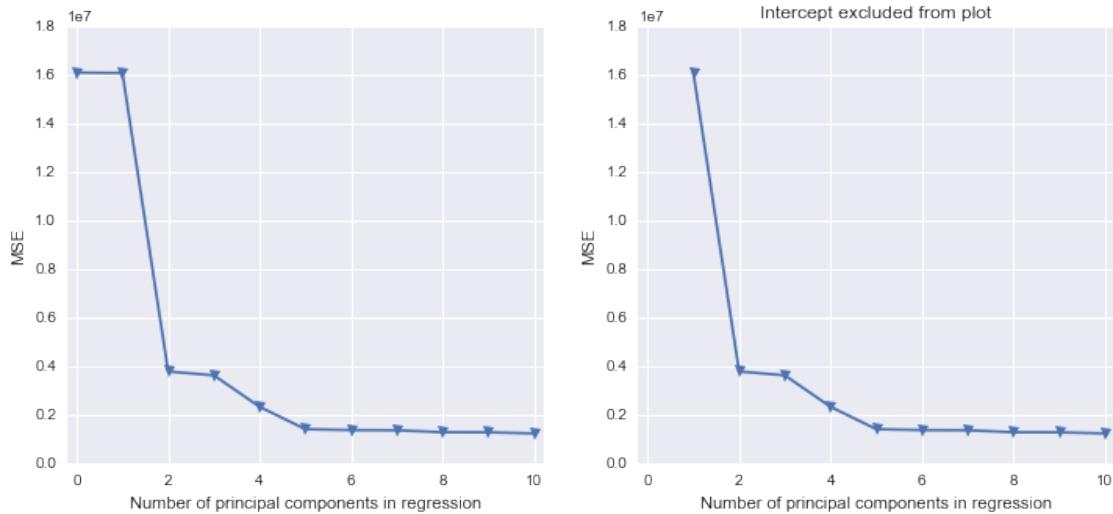
```
# Use 10 components adding one at a time
for i in np.arange(1,11):
    score = -1*cross_validation.cross_val_score(regr, X_reduced[:,i], Y_train.ravel(), cv=kf_
    mse.append(score)
```

```
fig, (ax1, ax2) = plt.subplots(1,2, figsize=(12,5))
ax1.plot(mse, '-v')
ax2.plot([1,2,3,4,5,6,7,8,9,10], mse[1:11], '-v')
ax2.set_title('Intercept excluded from plot')
```

```
for ax in fig.axes:
    ax.set_xlabel('Number of principal components in regression')
    ax.set_ylabel('MSE')
    ax.set_xlim((-0.2,10.2))
```

```
fig.show()
```

```
C:\Anaconda3\lib\site-packages\matplotlib\figure.py:397: UserWarning: matplotlib is currently using a non-
"matplotlib is currently using a non-GUI backend, "
```



```
In [23]: # Let's use 5 components, gives us the best score with the least number of components.
pcr_regr = linear_model.LinearRegression()
pcr_regr.fit(X_reduced[:, :5], Y_train)
```

```
score = mean_squared_error(Y_test, ridge_reg.predict(X_test))
print('Error on test data: {}'.format(score))
print('Value of M selected by CV: 5')
```

Error on test data: 1790804.8706922578

Value of M selected by CV: 5

6 Part F

Fit a PLS model on the training set, with M chosen by cross-validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
In [24]: from sklearn.cross_decomposition import PLSRegression
```

```
params = {'n_components': [2, 3, 4, 5, 7, 10]}
```

```
pls = PLSRegression()
pls_reg = GridSearchCV(pls, params, scoring='neg_mean_squared_error')
pls_reg.fit(X_train, Y_train)
```

```
score = mean_squared_error(Y_test, pls_reg.predict(X_test))
print("Error on test data: {}".format(score))
print("Value of M selected by CV: {}".format(pls_reg.best_params_['n_components']))
```

Error on test data: 10911489.731786955

Value of M selected by CV: 10

6.0.1 Part G

Comment on the results obtained. How accurately can we predict the number of college applications received? Is there much difference among the test errors resulting from these five approaches?

6.0.2 Answer

After trying all 5 approaches, they seemed to result in very similar scores on the test data except for PLS, which did much worse. The best we did was a score of about 1790788, which came from both Ridge Regression and PCR. This looks like the most accurate we can predict the number of college applications recieved. As stated earlier, all 5 models performed similarly save for PLS, which was much worse (about a factor of 5).

7 Problem 5

Generate data of the form $y = X\beta + \epsilon$, where X is a $n \times p$ matrix where $n = 51$, $p = 50$, and each $X_{ij} \sim N(0, 1)$. Also, generate the noise according to $\epsilon_i \sim N(0, \frac{1}{4})$. Let β be the all ones vector (for simplicity).

By repeatedly doing this experiment and generating fresh data (fresh X , y , and hence ϵ), but keeping β fixed, you will estimate many different solutions $\hat{\beta}$. Estimate the mean and variance of $\hat{\beta}$. Note that $\hat{\beta}$ is a vector, so for this exercise simply estimate the variance of a single component.

Choose regularization coefficients $\lambda = 0.01, 0.1, 1, 10, 100$ and repeat the above experiment. What do you observe? How do you explain this?

```
In [25]: def generate_data():
    X = np.random.randn(51, 50)
    epsilon = (1/4)*np.random.randn(51)
    beta = np.ones(50)
    y = np.dot(X, beta) + epsilon

    return (X, y)

def estimate_beta_hat(X, y, l=0):
    X_t = X.transpose()
    first_term = np.power(np.dot(X_t, X) + l*np.identity(50), -1)
    temp = np.dot(first_term, X_t)
    return np.dot(temp, y)

lambdas = [0, 0.01, 0.1, 1, 10, 100]
for l in lambdas:
    beta_hat_zeros = []
    for i in range(5000):
        X, y = generate_data()
        beta_hat = estimate_beta_hat(X, y, l=l)
        beta_hat_zeros.append(beta_hat[0])

    print('With Lambda = {}'.format(l))
    print('Mean: {}'.format(np.mean(beta_hat_zeros)))
    print('Variance: {}'.format(np.var(beta_hat_zeros)))
    print()
```

```
With Lambda = 0
Mean: -557.0247249562078
Variance: 2041609410.5309222
```

```
With Lambda = 0.01
Mean: 582.5252144966123
Variance: 1132450700.985211
```

```
With Lambda = 0.1
Mean: 133.53251112850842
```

Variance: 237240137.79508066

With Lambda = 1

Mean: -1711.9628437159895

Variance: 16216570457.940891

With Lambda = 10

Mean: -1295.1044065392912

Variance: 5935390658.723576

With Lambda = 100

Mean: 4939.4450696976055

Variance: 68305244969.84221

7.0.1 Answer

When introducing a higher regularization coefficient, the mean and variance of the estimated beta value should both shrink. Because a higher value of λ adds more penalty to higher beta values, the beta values should naturally tend toward zero. Especially in our example, since our true betas are all ones, this should help. The mean should shrink as well as the variance.